



reveals



SevenBridges

DNA Sequence Assembly

Milan Kovacevic
milan.kovacevic@velsera.com

Agenda

- Assembly – basic overview
- DeBruijnGraph assembly
- String Graphs in assembly
- Assembly metrics
- Python examples

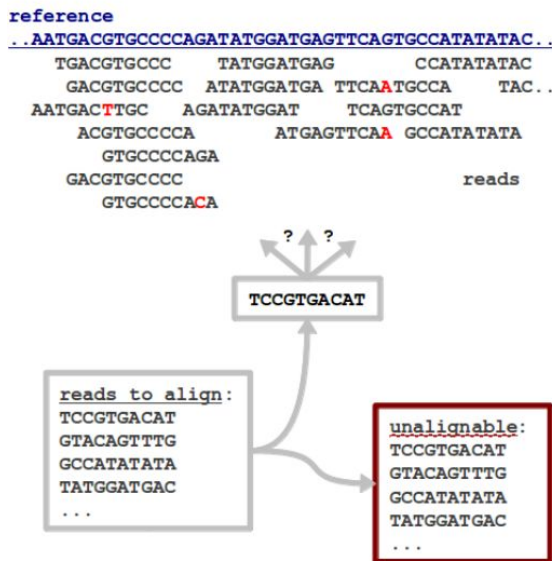
Recap

- What is PCR?
- What is a flow cell?
- What are paired-end reads?
- What to do with the sequencing reads?
- Human Genome Project?
- How many chromosomes do people have?
- Define the Central Dogma, genes and exome.
- What do we store in FASTA, FAI and FASTQ file formats?

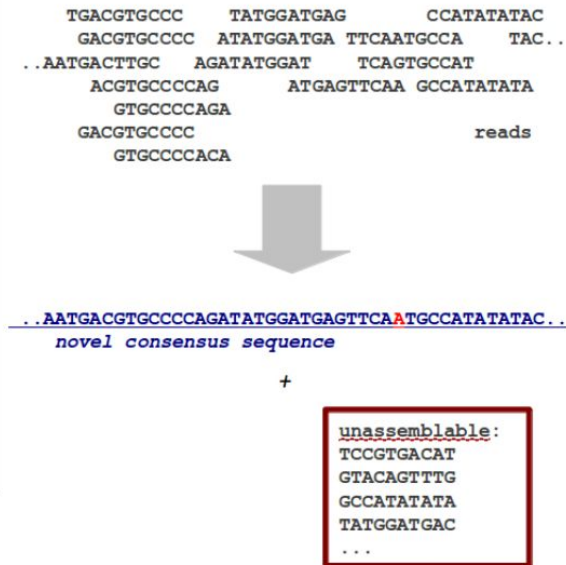
De novo assembly

- **Assembly:** Set of sequences which best approximate the original sequenced material
- Why?
- Reference
- Novel features

Alignment



Assembly



De novo assembly

- Best case scenario: check every pair of (long) reads for overlaps
- Computationally expensive: n reads, $\sim n^2$ operations
- Our task: how to get the **best assemblies** for the **smallest expense** – in terms of sequencing and computational expenses

Read1 - TTTGGTGCTCTTCGAAAAGGGATCTTCGAGAGAGATCTCGCGATAAGGTTG

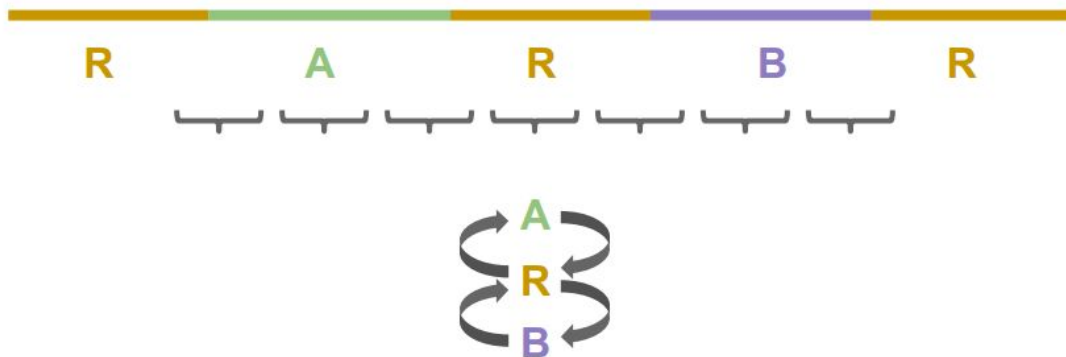
Read2 - GAGAGAGATCTCGCGATAAGGTTGAAGTAGAAAAATGTGTGTGGTGAA

overlap

TTTGGTGCTCTTCGAAAAGGGATCTTCGAGAGAGATCTCGCGATAAGGTTG
GAGAGAGATCTCGCGATAAGGTTGAAGTAGAAAAATGTGTGTGGTGAA

De novo assembly

- “[...] **repeats** are the single biggest impediment to all assembly algorithms and sequencing technologies” - Koren 2012 Nature Biotechnology
- Short read libraries have a hard time resolving large repeats, even with mate-pair, or "jumping" libraries



A typical assembly pipeline

Raw reads



Reads error correction

Error corrector



Single reads indexing

Assembler

Single reads assembly



Paired reads scaffolding

Scaffolder



Scaffolds gap-closing

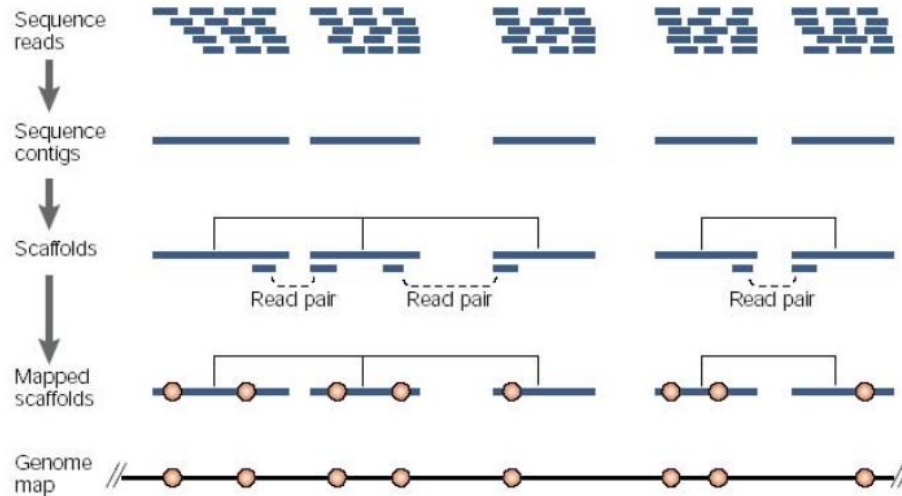
Gapcloser



Contigs, scaffolds

A typical assembly pipeline

de novo whole-genome shotgun assembly



Vocabulary

- **Paired end reads:** read1, insert < 500 bp, read2
- **Mate pair reads:** read1, insert > 1 kbp, read2
- **k-mer:** any sequence of length k
- **Contig:** gap-less assembled sequence
- **Scaffold:** sequence which may contain gaps

Short-Insert Paired End Reads

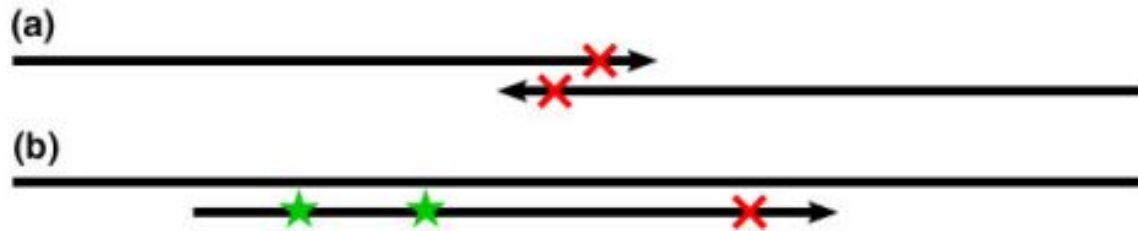


Long-Insert Paired End Reads (Mate Pair)




Error correction

- Assembly: errors (mostly at 3' read ends) disrupt overlaps
- Alignment: where real variation is present, errors add to differences and disrupt alignment



Error correction

- Determining trusted and untrusted k-mers
- Untrusted -> trusted, using highest likelihood



e.g.:

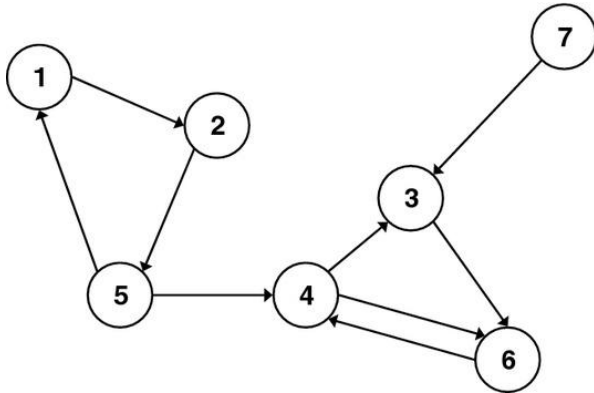
CCGGATTTACGCCGAGTGA x 43

CCGGATTTACGCTGAGTGA x 1

^

Graphs

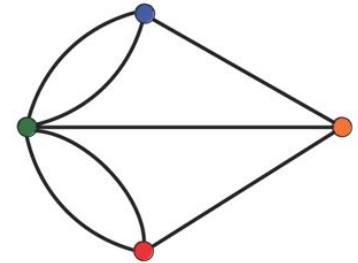
- A **graph** is a set of nodes and a set of edges
- Edges can be **directed** or **not**



a



b



Graphs and de novo assembly

- **Overlaps** between reads – fundamental information used for de novo assembly
- **Graphs** enable representation of these overlaps
- Two different types of graphs for sequencing data are known :
 - **de Bruijn** graphs
 - **string** graphs



De Bruijn graphs

- For a fixed integer **k**:
 - **nodes** - all **k-1**-mers present in reads
 - **edges** - for each **k**-mer **x** present in reads there is an edge between **k-1**-mer **prefix** of **x**, and **k-1**-mer **suffix** of **x**
- Example for a single read and $k = 4$:

AAC**TG**



De Bruijn graphs

- In assembly, we always identify a read with its reverse complement
- **k** is practically exclusively an **odd** integer as a result (otherwise that causes ambiguities in the strand-specificness of the graph)



De Bruijn graphs

Example for multiple reads and $k = 4$:


AACTG
ACTGC
CTGCT



De Bruijn graphs

Similar example for multiple reads and $k = 4$:

ACTG
TGCT

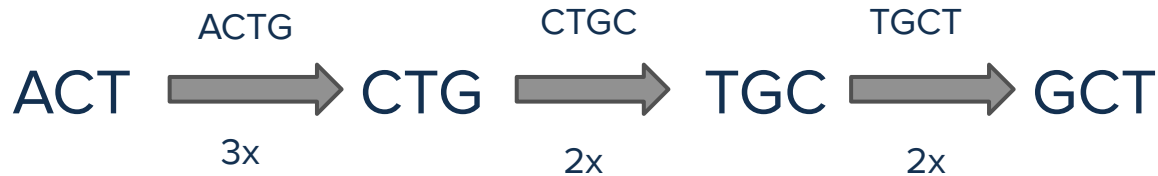


No connection between two graphs because there are no overlaps $\geq k-1$

De Bruijn graphs

- What happens if we add redundancy?
- Previous example for multiple reads and $k = 4$:

ACTG
ACTG
ACTG
CTGC
CTGC
TGCT
TGCT



De Bruijn graphs

- How does a sequencing error impact de Bruijn graph?
- Previous example for multiple reads and $k = 4$:

ACTG

ACTG

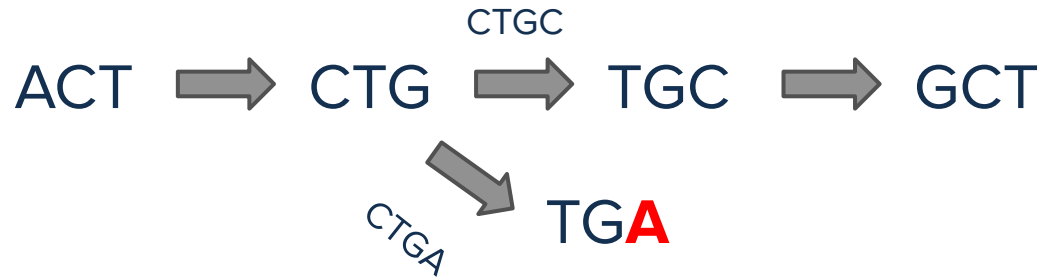
ACTG

CTGC

CTGA

TGCT

TGCT




De Bruijn graphs

- What is the effect of a small repeat on the graph?
- Example for multiple reads and $k = 4$:

ACTG
CTGC
TGCT
GCTG
CTGA
TGAC

ACTGCTGAC



Summary

- Given any sequence and k-mer size, we can create a De Bruijn graph in a unique manner.
- The other direction is not true. All De Bruijn graphs cannot be resolved into unique sequences. Unless the De Bruijn graph is in its simplest form, it usually resolves into many possible sequences.
- Larger the k-mers, more easy it is to convert the De Bruijn graph into a unique sequence.

Summary

What are the limitations of De Bruijn graphs?

- Reads are immediately split into shorter k-mers; can't resolve repeats as well as overlap graph
- Read coherence is lost. Some paths through De Bruijn graph are inconsistent with respect to input reads.
- Single most important benefit of De Bruijn graph is speed and simplicity.

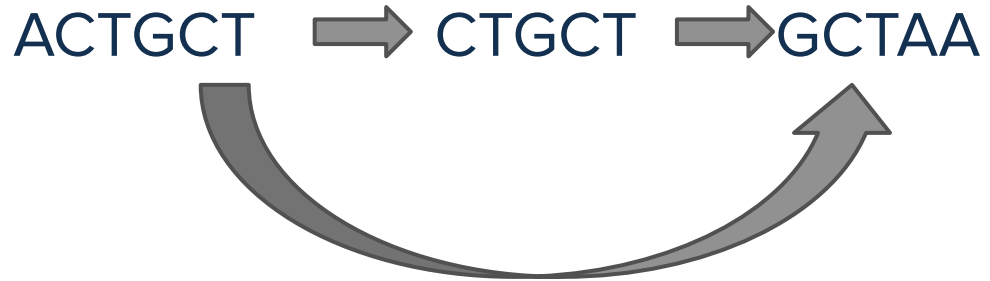
String graphs

- String graph is a modification of an overlap graph.
- Overlap graph:
 - **nodes** – reads
 - **edges** – two nodes are connected if they overlap
- What does “overlap” mean? Similar to “align”, no silver bullet

Overlap graph

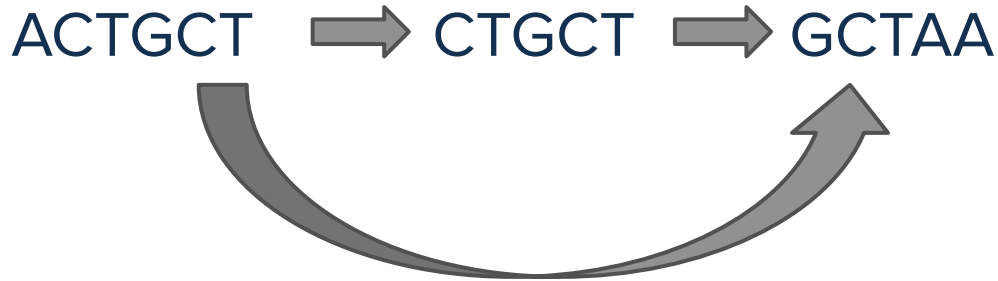
- Given $p > 0$, we say that r_1 and r_2 overlap if a suffix of r_1 of **length** $\geq p$ is exactly a prefix of r_2 of same length
- Example for multiple reads and $p = 3$:

ACTGCT
CTGCT
GCTAA



String graph

- A string graph is obtained from an overlap graph by removing redundancy:
 - **redundant reads** (those fully contained in another read)
 - **transitively redundant edges** (if $x \rightarrow y$ and $y \rightarrow z$, then keep only $x \rightarrow z$)
- Overlap and string graph for previous example:



String vs De Bruijn graph

String and De Bruijn graph for the same example:

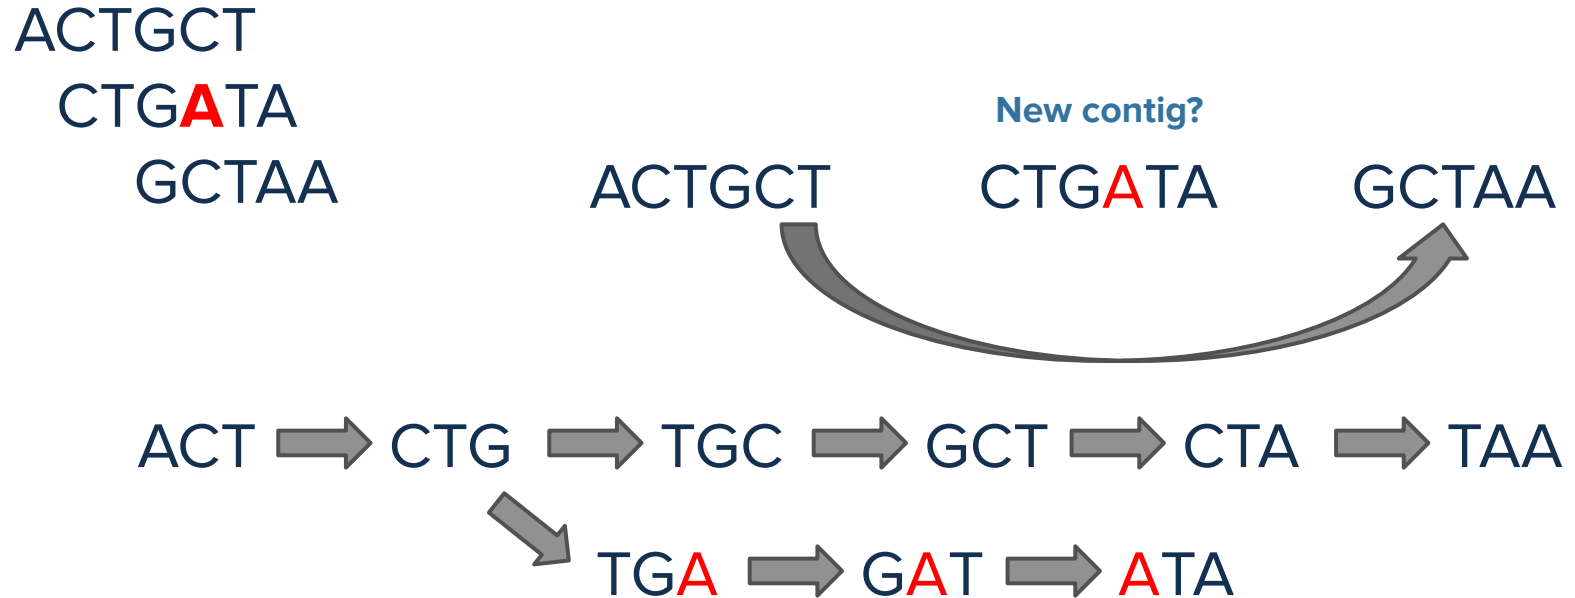
ACTGCT
CTGCTA
GCTAA

ACTGCT → CTGCTA → GCTAA

ACT → CTG → TGC → GCT → CTA → TAA

String vs De Bruijn graph

Previous example, with error:

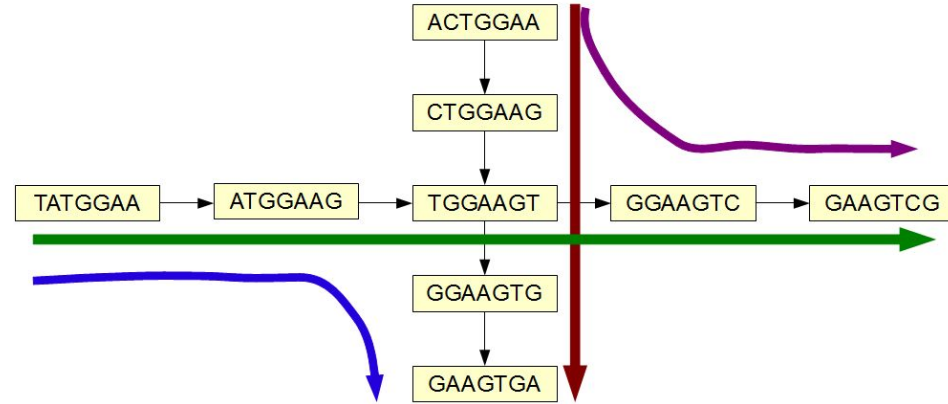


String vs De Bruijn graph

- Which is better?
 - String graphs capture whole read information
 - de Bruijn graphs are conceptually simpler:
 - single node length
 - single overlap definition
- Historically, **string graphs** were used for **long reads**
- and **De Bruijn** graphs for **short reads**

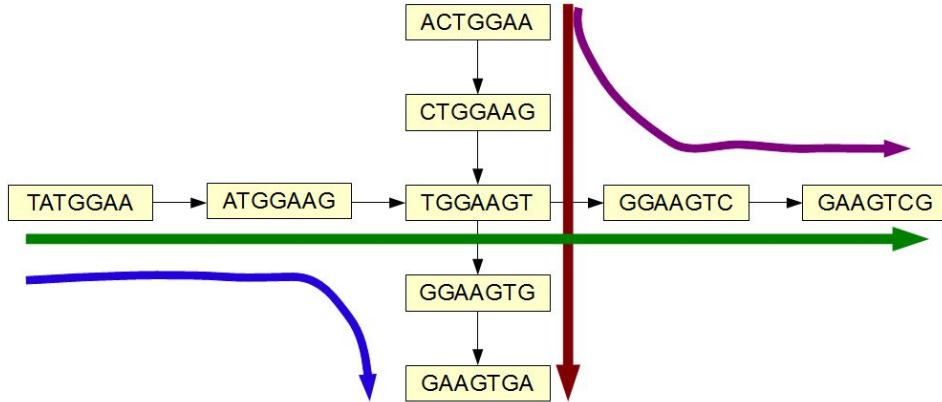
De Bruijn graph

- How do we choose **k**?
- Large **k**:
 - sequencing errors bigger problem
 - graph less connected
- Small **k**:
 - ambiguous paths
- Balance, try different values for **k**



Original sequences – **TATGGAAGTCG**, **ACTGGAAGTGA**

De Bruijn graph



Original sequences – TATGGAAGTCG, ACTGGAAGTGA

k=8

ACTGGAAGT
↓
CTGGAAGTG
↓
TGGAAGTGA

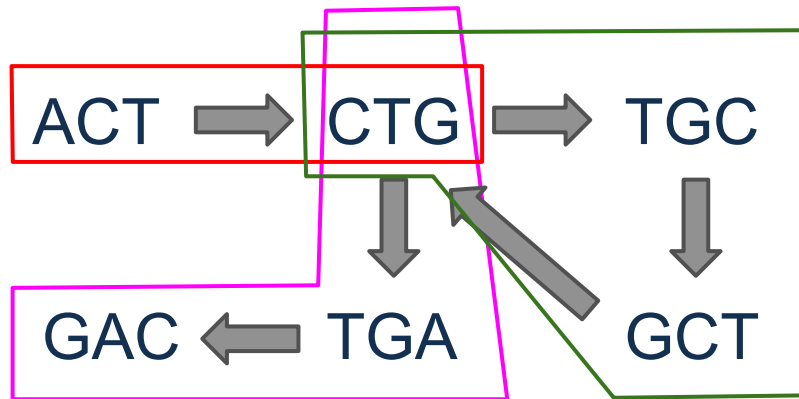
TATGGAAGT
↓
ATGGAAGTC
↓
TGGAAGTCG

k=10

Contigs construction

Contigs construction in practice: return a set of paths covering the graph, such that all possible assemblies contain these paths

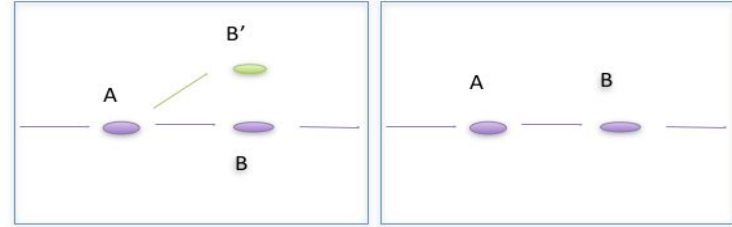
ACTG
CTGCT
CTGAC



Graph topology based error correction

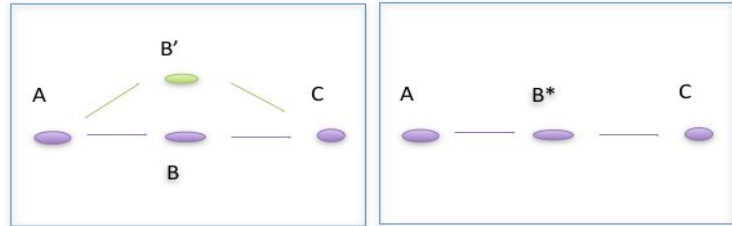
—Errors at end of read

- Trim off 'dead-end' tips



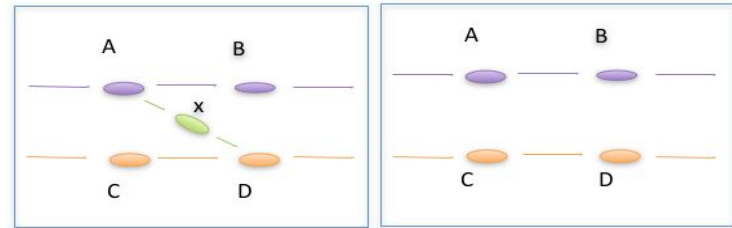
—Errors in middle of read

- Pop Bubbles



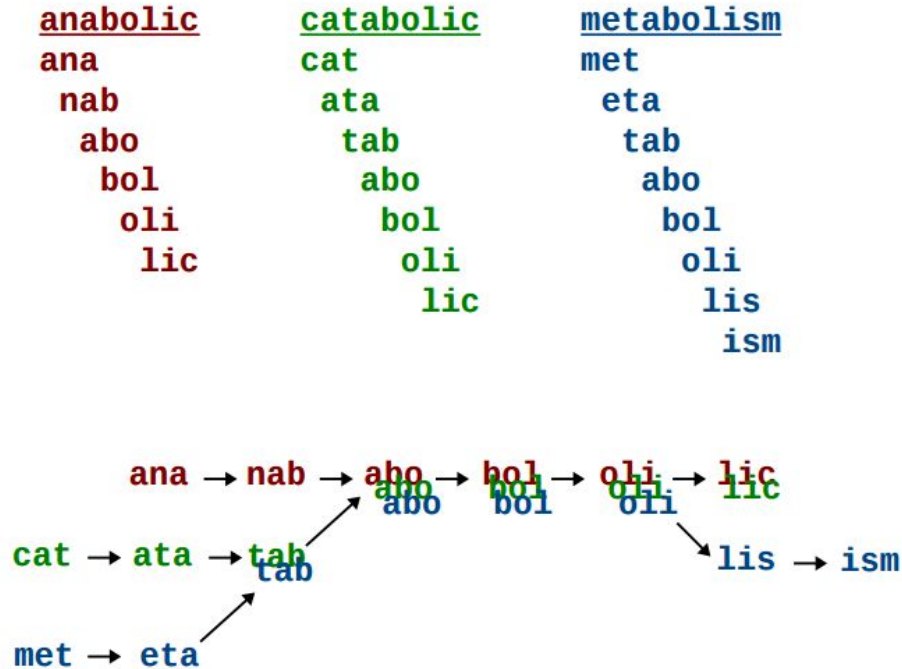
—Chimeric Edges

- Clip short, low coverage nodes



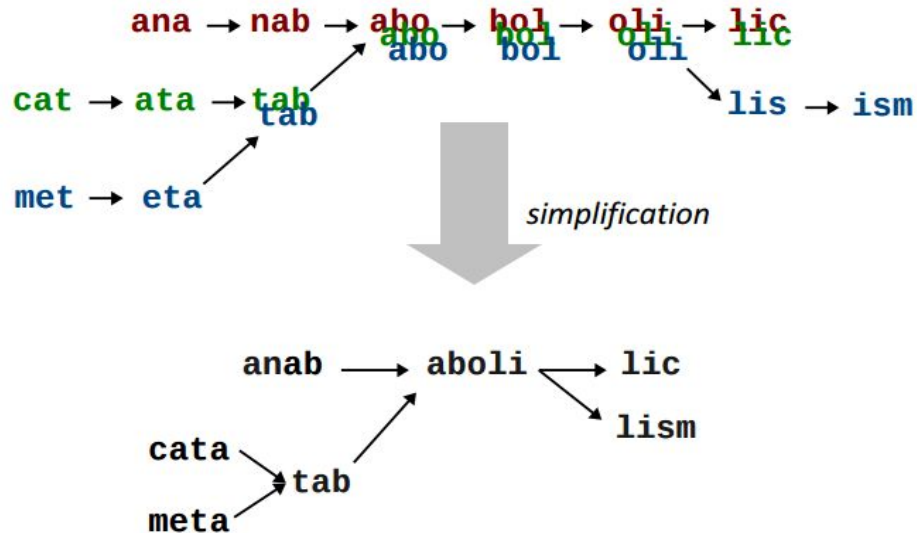
Building De Bruijn graph

K-mer size equal to 4

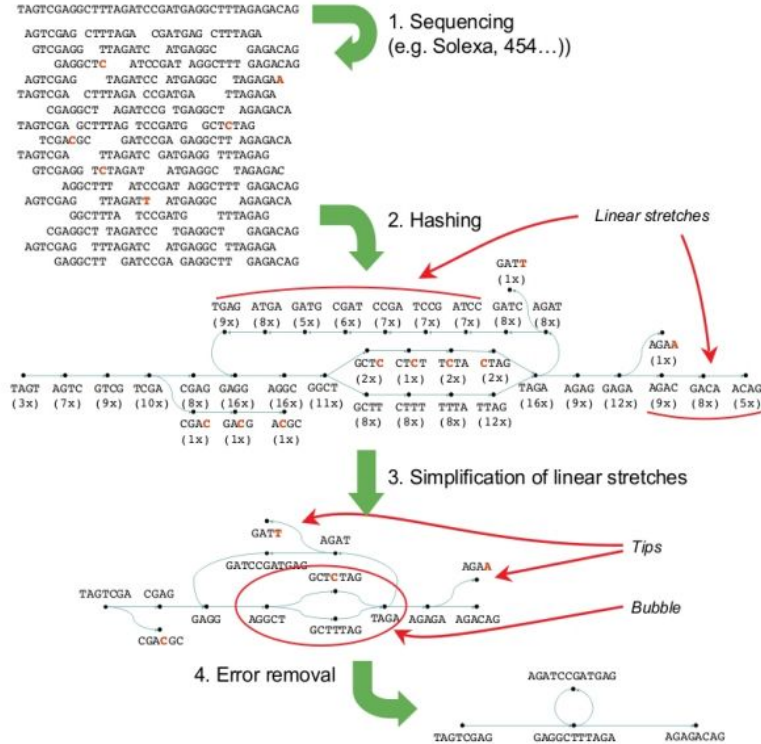


Building De Bruijn graph

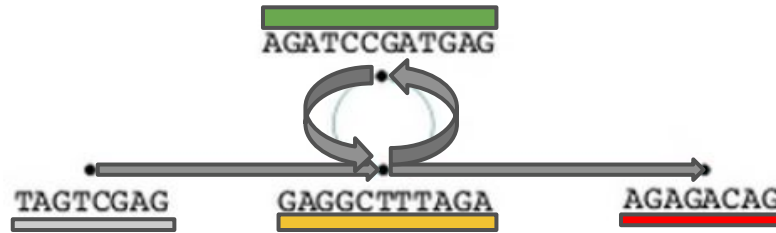
Weight = number of times k-mer occurs



Contigs construction



Contigs construction



Potential assemblies:

TAGTCGAGGAGGCTTTAGAAAGATCCGATGAGGAGGCTTTAGAAAGAGACAG

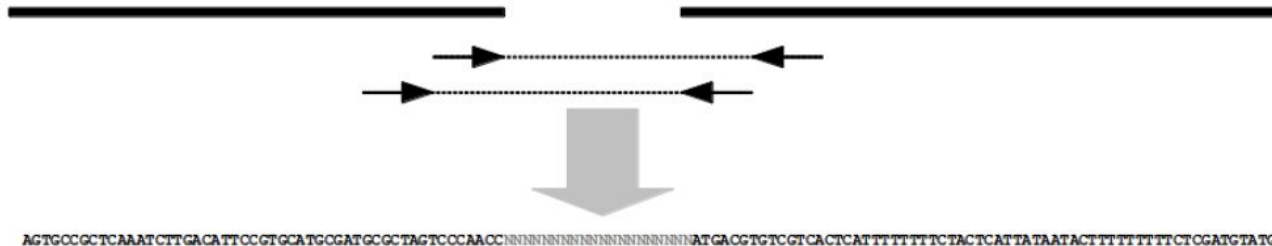
TAGTCGAGGAGGCTTTAGAAAGATCCGATGAGGAGGCTTTAGAAAGATCCGATGAGGAGGCTTTAGAAAG

AGACAG

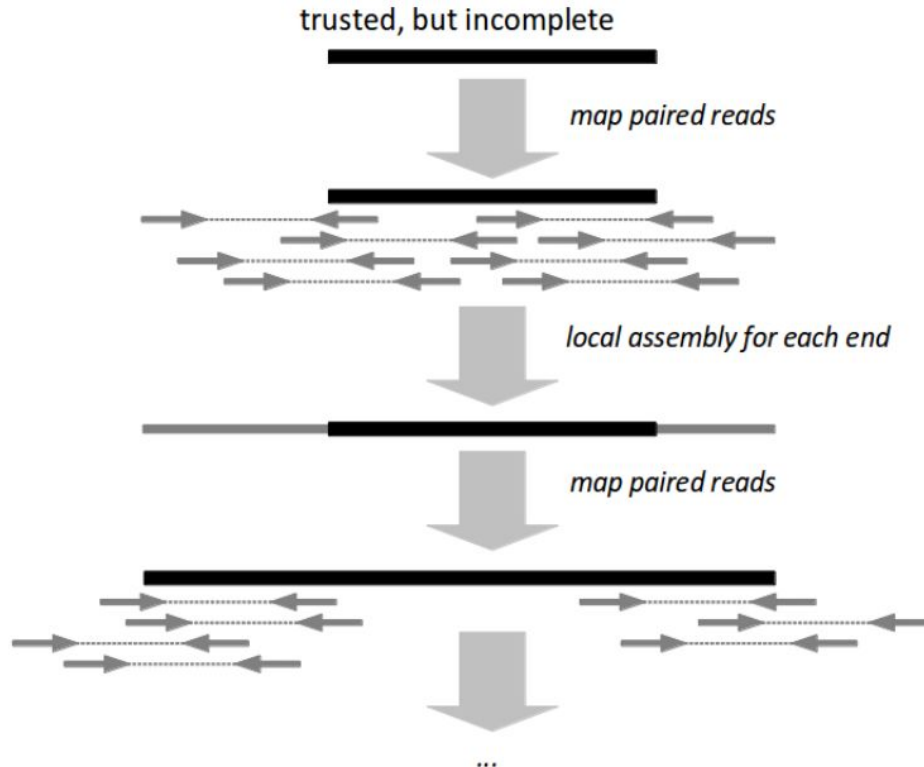
...

Scaffolds construction

- Scaffolding using pairing information
- This is why libraries containing **multiple insert sizes** are used:
 - Gnerre 2011:
 - 45x overlapping PE reads (insert size: 180 bp, reads: >100 bp)
 - 45x short jump MP reads (insert size: 3 kb)
 - 5x (optional) long jump MP reads (insert size: 6 kb)
 - 1x (optional) fosmid jump MP reads (insert size: 40 kb)
 - Ribeiro 2012:
 - 50x overlapping PE reads (insert size: 180 bp, reads: >100 bp)
 - 50x PacBio reads (reads: 1-3 kb)
 - 50x long jump MP reads (insert size: 2-10 kb)

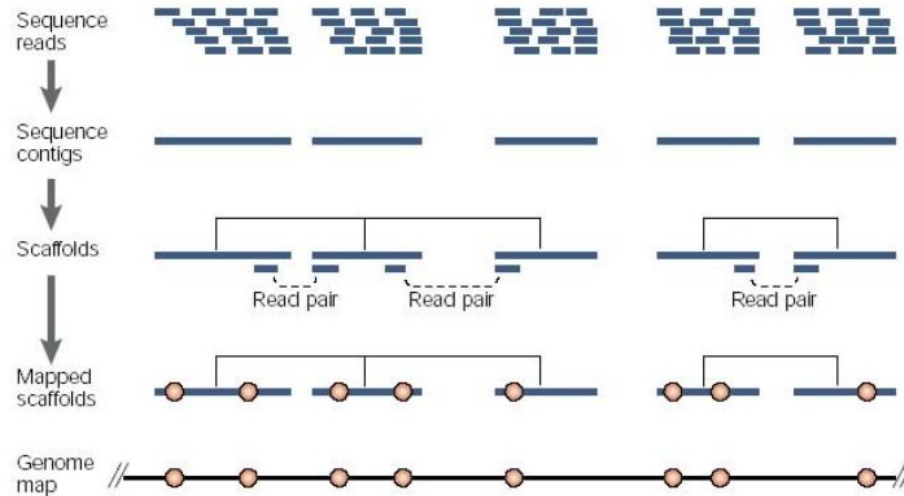


Gap closing / contig extension



Summary

de novo whole-genome shotgun assembly




Metrics

- How do we evaluate different assemblies?
- There is no trivial ranking between assemblies
- Reference-free metrics:
 - Number of contigs / scaffolds
 - Total length of the assembly
 - Length of the largest contig / scaffold
 - Percentage of gaps in scaffolds
 - **Nx, NGx of contigs / scaffolds**
 - **Internal consistency**
 - Number of predicted genes
- Reference-using metrics:
 - **Coverage**
 - Assembly errors

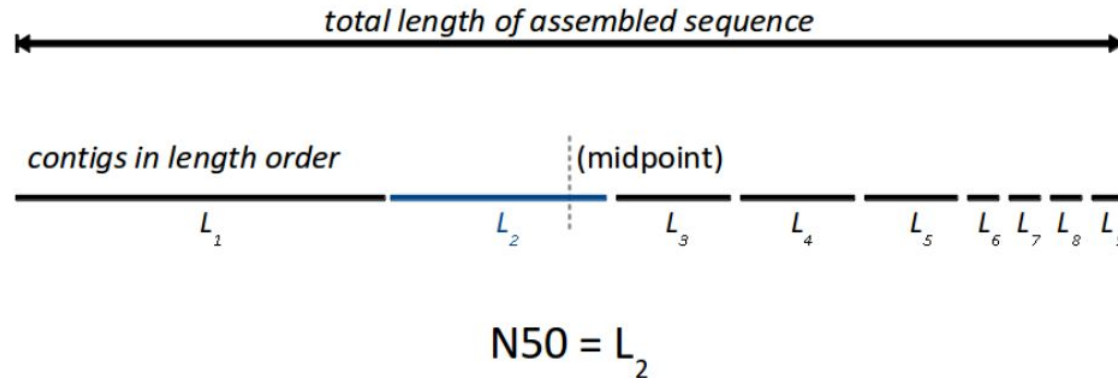
Nx / NGx of contigs / scaffolds

- Nx, where x is percentage value
- **N50** is analogous to **median**
- Example set {7, 3, 1, 1, 1}
- Median is 1
- N50 is 7



Nx / NGx of contigs / scaffolds

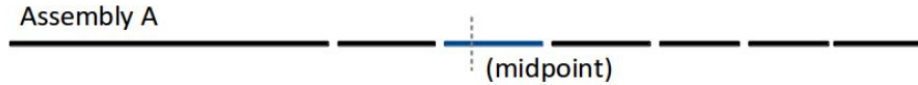
- **Nx** is largest contig length at which longer contigs cover x% of the total assembly length




- A practical way to compute Nx:
 - Sort contigs by decreasing lengths
 - Take the first contig (the largest): does it cover x% of the assembly?
 - If yes, this is the Nx value. Else, repeat by trying the next one (the second largest)

Nx / NGx of contigs / scaffolds

What's the problem with Nx?




Assembly B's N50 >> Assembly A's N50 (??)



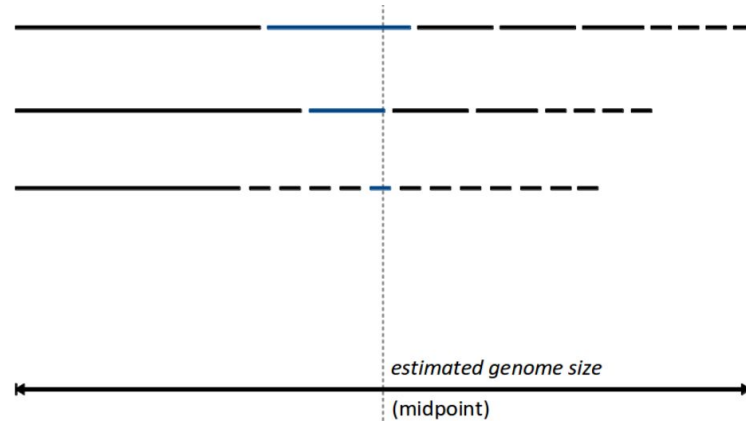
Nx / NGx of contigs / scaffolds

Solution - NGx



Nx / NGx of contigs / scaffolds

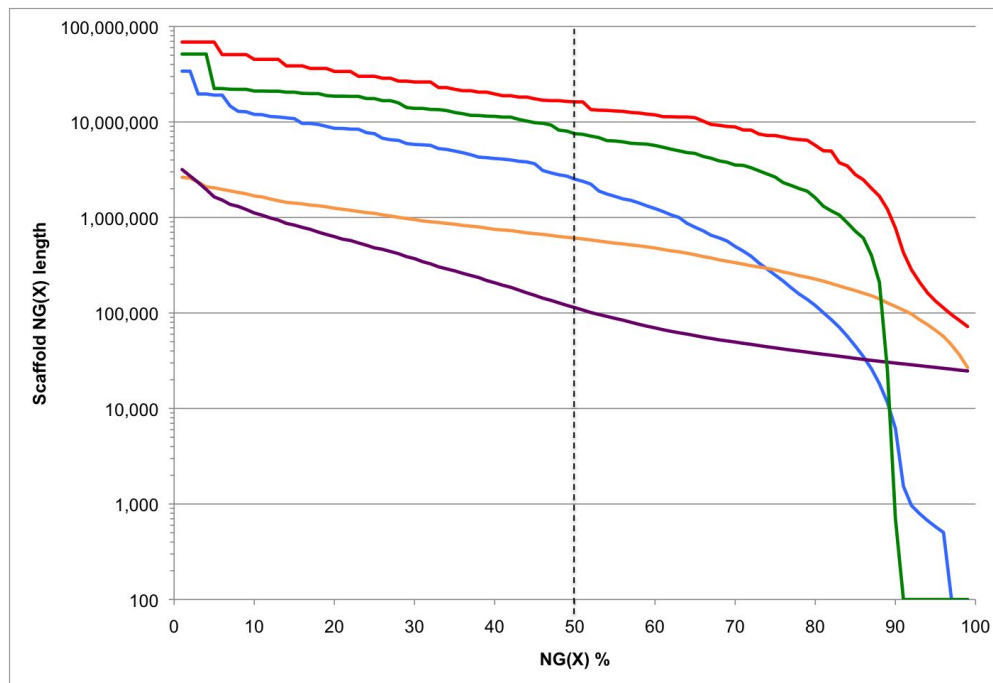
- **NGx** is largest contig length at which longer contigs cover x% of the total genome length



- A practical way to compute NGx:
 - Sort contigs by decreasing lengths
 - Take the first contig (the largest): does it cover x% of the genome?
 - If yes, this is the NGx value. Else, repeat by trying the next one (the second largest)


Nx / NGx of contigs / scaffolds

Assemblathon 2



Metrics

- **Internal consistency:** percentage of paired reads correctly aligned back to the assembly (happy pairs)
- **Coverage:** percentage of bases in the reference which are covered by the alignment



Further reading

- A short and nice paper containing essential explanations about assembly graphs - How to apply de Bruijn graphs to genome assembly, <http://www.nature.com/nbt/journal/v29/n11/pdf/nbt.2023.pdf>
- Assembly lecture from MITs “Foundations of Computational and Systems Biology”: <https://www.youtube.com/watch?v=ZYW2AeDE6wU>
- Assemblathon - assembly (<http://assemblathon.org/>) competition and the paper: Assemblathon 2: evaluating de novo methods of genome assembly in three vertebrate species
- Nice assembly ppt presentation - http://www.cs.jhu.edu/~langmea/resources/lecture_notes/assembly_dbg.pdf
- Nice assembly ppt presentation from MIT: https://ocw.mit.edu/courses/biology/7-91j-foundations-of-computational-and-systems-biology-spring-2014/lecture-slides/MIT7_91JS14_Lecture6.pdf

Exercise 1

Write a function “kmerize” which takes in a string (read) and a k-mer length, and returns a list of all unique k-mers present in the string.

Run example:

```
> kmerize("ACGCGTCGC", 3)
> ['ACG', 'CGC', 'GCG', 'CGT', 'GTC', 'TCG']
```

Exercise 1b

Write a function “kmer_count” which takes in a string (read) and a k-mer length, and returns a dictionary containing all unique kmers and number of appearances each of these kmers.

Run example:

```
> kmer_count("AAACCCGCGC", 2)
> {'AA': 2, 'AC': 1, 'CC': 2, 'CG': 2, 'GC': 2}
```

Exercise 2

Write a function “de_bruijn_ize” which takes in a string (read) and a k-mer length, and returns a tuple with all unique k-1-mers (nodes) and a list holding, for each k-mer, its left k-1-mer and its right k-1-mer in a pair (edges).

Run example:

```
> de_bruijn_ize("ACGCGTCG", 3)
> ({'AC', 'GT', 'GC', 'CG', 'TC'},
  [('AC', 'CG'), ('CG', 'GC'), ('GC', 'CG'), ('CG', 'GT'),
   ('GT', 'TC'), ('TC', 'CG')])
```

Exercise 3

Simple De Bruijn Graph implementation with Eulerian walk-finder
Solution (page 16):

http://www.cs.jhu.edu/~langmea/resources/lecture_notes/assembly_dbg.pdf

Code:

http://nbviewer.jupyter.org/github/BenLangmead/comp-genomics-class/blob/master/notebooks/CG_deBruijn.ipynb